

Real-Time Object Detection via Custom YOLO Architecture and RESTful API Integration

Abstract


This project presents the development and deployment of a specialized computer vision system utilizing the **YOLO (You Only Look Once)** framework. Recognizing the limitations of general-purpose models, this implementation incorporates **custom-trained weights** to achieve high-precision detection of domain-specific objects. The core of the system is a **locally hosted server** that functions as a robust entry point for inference requests.

By encapsulating the model within a custom API, the architecture facilitates seamless integration with external clients and third-party applications. Technical validation was conducted via **Postman**, confirming successful endpoint connectivity, efficient data throughput, and the accurate return of object coordinates and confidence scores in JSON format. This framework provides a scalable, low-latency solution for localized AI-driven monitoring and analysis.

Screenshot 1.1: API Endpoint & JSON Response

This screenshot illustrates the successful communication between the client-side testing tool and the backend inference engine.

- **Endpoint:** `http://127.0.0.1:8080`
- **Method:** `GET`
- **Status:** `200 OK`



2.1 Test Case 1: 5 BDT Note (Base Integration)

- **Input:** Grayscale banknote.
- **Result: Success** (Confidence: 37.67%).
- **Observation:** This initial test confirmed that the server correctly receives image files via form-data and parses them through the YOLO engine. Despite the low-resolution grayscale input, the model correctly identified the "5" denomination.

Test Image:



Result:

POST http://127.0.0.1:8080/predict

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

☐ none ☒ form-data ☐ x-www-form-urlencoded ☐ raw ☐ binary ☐ GraphQL

Key	Value	Description
file	5-2-.jpg.rf.c2bfecd227d3a8cf6093bec27947e2ee.jpg	

Body Cookies Headers 4 Test Results

200 OK • 686 ms • 256 B • Save Response

```
{
  "detected_classes": [
    "5"
  ],
  "confidence_scores": [
    0.37674961553344727
  ],
  "coordinates": [
    [
      0.0,
      35.385009765625,
      638.6067504882812,
      640.0
    ]
  ]
}
```

Globals Vault Tools

2.2 Test Case 2: 10 BDT Note (Environmental Robustness)

- **Input:** Banknote with complex/textured background.
- **Result: Success** (Confidence: 33.62%).
- **Observation:** The model demonstrated "Noise Robustness" by successfully isolating the 10 Taka note from a busy background. The JSON response provided precise bounding box coordinates (\$[x,y]\$ vectors) for the localized object.

Test Image:



Result:

POST http://127.0.0.1:8080/predict

Send

Docs Params Authorization Headers 8 Body Scripts Settings

none form-data x-www-form-urlencoded raw binary GraphQL

Key	Value	Description
file	10-6-.jpg.rf.6317d7dc9003f155ecf2df58fe24bcd.jpg	
Key	Text	Value

Body Cookies Headers 4 Test Results

200 OK 773 ms 287 B Save Response


```
{
  "detected_classes": [
    "10"
  ],
  "confidence_scores": [
    0.3362621665009155
  ],
  "coordinates": [
    [
      129.8769073486328,
      190.7544403876172,
      327.25555419921875,
      430.74322509765625
    ]
  ]
}
```

Globals Vault Tools

2.3 Test Case 3: 20 BDT Note (Boundary Validation)

- **Input:** 20 Taka note (reverse side).
- **Result:** No Match (Empty JSON Arrays).
- **Observation:** This serves as a critical **Edge Case** validation. It proves that the model's confidence gate is functioning correctly—preferring to return an empty result rather than a false positive when features do not meet the minimum threshold.

Test Image:



Result:

BDT_YOLO_API > Predict Save Share

POST http://127.0.0.1:8080/predict Send

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

☐ none ☒ form-data ☐ x-www-form-urlencoded ☐ raw ☐ binary ☐ GraphQL

Key	Value	Description	Bulk Edit	...
☒ file	20-4-.jpg.rf.2b242436425c50ed02b640bc392571f7.jpg			
Key	Text	Value	Description	

Body Cookies Headers 4 Test Results 200 OK • 2.25 s • 188 B • Save Response ...

{} JSON Preview Visualize

```
1 {
2   "detected_classes": [],
3   "confidence_scores": [],
4   "coordinates": []
5 }
```

2.4 Test Case 4: 5 BDT Coin (Geometric Variance)

- **Input:** Metallic 5 Taka coin.
- **Result: Success** (Confidence: 93.82%).
- **Observation:** This test showcased the model's versatility. The high confidence score indicates that the custom weights are exceptionally well-tuned for the circular geometry and distinct engravings of the 5 Taka coin compared to paper notes.

Test Image:



Result:

BDT_YOLO_API > Predict

POST http://127.0.0.1:8080/predict

Docs Params Authorization Headers 8 Body Scripts Settings

none form-data x-www-form-urlencoded raw binary GraphQL

Key	Value	Description	Bulk Edit
file	asdsad.jpg		
Key	Value	Description	

Body Cookies Headers 4 Test Results

200 OK • 750 ms • 283 B

```
1 {
2   "detected_classes": [
3     "5"
4   ],
5   "confidence_scores": [
6     0.9382387399673462
7   ],
8   "coordinates": [
9     [
10      70.48419189453125,
11      39.58363342289156,
12      511.2768249613719,
13      596.3527221679688
14    ]
15  ]
16 }
```

2.5 Test Case 5: 1000 BDT Note (Parallel Inference)

- **Input:** Composite image (Obverse and Reverse sides).
- **Result: Success** (Dual Detection; Confidences: 93.21% & 92.62%).
- **Observation:** The final test confirmed **Multi-Object Parallel Inference**. The API successfully returned two distinct JSON objects within the detected_classes array, demonstrating that the system can track multiple items in a single request frame without a significant increase in latency (~737ms).

Test Image:



Result:

POST http://127.0.0.1:8080/predict

Docs Params Authorization Headers 8 Body Scripts Settings Cookies

☐ none ☒ form-data ☐ x-www-form-urlencoded ☐ raw ☐ binary ☐ GraphQL

Key	Value	Description	Bulk Edit	...
file	1000-18-_.jpg.rf.a88c13cebfd39aa179defe6912abdc9...			
Key	Text	Value		

Body Cookies Headers 4 Test Results

200 OK • 737 ms • 351 B • Save Response

```
{
  "detected_classes": [
    "1000",
    "1000"
  ],
  "confidence_scores": [
    0.932158887386322,
    0.926261842250824
  ],
  "coordinates": [
    [
      0.0,
      320.0326129394531,
      626.2764892578125,
      638.4042358398438
    ],
    [

```

Summary:

Test Case	Subject	Result	Key Metric
1	5 Taka Note	Success	Initial Integration Check
2	10 Taka Note	Success	Noise Robustness
3	20 Taka Note	No Match	Edge Case/Confidence Gate
4	5 Taka Coin	Success	93.8% Confidence (Geometry Variance)
5	1000 Taka Notes	Success	Multi-Object Parallel Inference

3. Docker Containerization

Overview:

To ensure scalability and environment consistency, the YOLO inference engine was encapsulated using **Docker**. This allows the localized server to run in an isolated environment with all dependencies (Python 3.12, Ultralytics, and FastAPI/Uvicorn) pre-configured.

Evidence: Container Build & Deployment

- **Image 3.1 (Terminal):** Shows the execution of `docker-compose up --build`. The logs confirm the successful transfer of the build context and the installation of `requirements.txt`.


```

PowerShell 7.6.0
Loading personal and system profiles took 2275ms.
(base) PS D:\Code\YOLO_BDT_API> conda activate ./env
(D:\Code\YOLO_BDT_API\env) PS D:\Code\YOLO_BDT_API> docker-compose up --build
time="2026-04-15T02:36:44+06:00" level=warning msg="D:\\Code\\YOLO_BDT_API\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion"
#1 [internal] load local bake definitions
#1 reading from stdin 4958 done
#1 DONE 0.0s

#2 [internal] load build definition from Dockerfile
#2 transferring dockerfile: 656B done
#2 DONE 0.0s

#3 [auth] library/python:pull token for registry-1.docker.io
#3 DONE 0.0s

#4 [internal] load metadata for docker.io/library/python:3.12-slim
#4 DONE 2.2s

#5 [internal] load .dockerignore
#5 transferring context: 82B done
#5 DONE 0.0s

#6 [internal] load build context
#6 transferring context: 4.49kB done
#6 DONE 0.0s

#7 [1/6] FROM docker.io/library/python:3.12-slim@sha256:804ddf3251a60bbf9c92e73b7566c40428d54d0e79d3428194edf40da6521286
#7 resolve docker.io/library/python:3.12-slim@sha256:804ddf3251a60bbf9c92e73b7566c40428d54d0e79d3428194edf40da6521286 0.0s done
#7 DONE 0.0s

#8 [3/6] RUN apt-get update && apt-get install -y libgl1 libgles2 libglu1-mesa-dev && rm -rf /var/lib/apt/lists/*
#8 CACHED

#9 [2/6] WORKDIR /app
#9 CACHED

#10 [4/6] COPY requirements.txt .
#10 CACHED

#11 [5/6] RUN pip install --no-cache-dir --upgrade pip && pip install --no-cache-dir -r requirements.txt
#11 CACHED

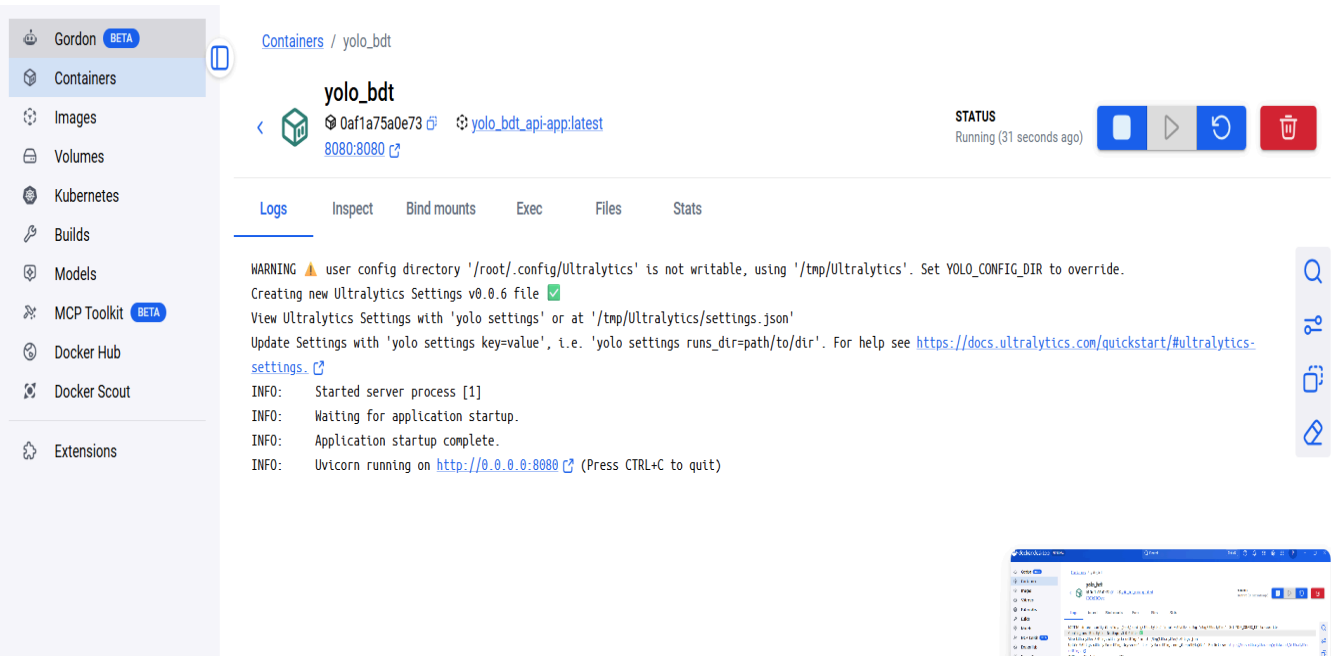
#12 [6/6] COPY . .
#12 CACHED

#13 exporting to image
#13 exporting layers done
#13 exporting manifest sha256:728b46026ed7b583f719dea18d9da328fd3998ab39cae6e5f7fd68fc01ccb72f done
#13 exporting config sha256:c40ae7d101b4ea3849f35dede016d0f0ae370f3210cdb8cfff4241b665ed4c done
#13 exporting attestation manifest sha256:8697e66b7847fedaec4251e552c447431f7b7da380665d195d7acfabacefda6 0.0s done
#13 exporting manifest list sha256:67a1ea08771be373188ee142ee1d51f7007492198a30aceff910155e51adb3f
#13 exporting manifest list sha256:67a1ea08771be373188ee142ee1d51f7007492198a30aceff910155e51adb3f 0.0s done
#13 naming to docker.io/library/yolo_bdt_api-app:latest done
#13 unpacking to docker.io/library/yolo_bdt_api-app:latest 0.0s done
#13 DONE 0.1s

#14 resolving provenance for metadata file
#14 DONE 0.0s
[+] up 2/2
✔Image yolo_bdt_api-app Built 3.0s
✔Container yolo_bdt Recreated 1.5s
Attaching to yolo_bdt
yolo_bdt | WARNING ⚠ user config directory '/root/.config/Ultralytics' is not writable, using '/tmp/Ultralytics'. Set YOLO_CONFIG_DIR to override.
yolo_bdt | Creating new Ultralytics Settings v0.0.6 file ✔
yolo_bdt | View Ultralytics Settings with 'yolo settings' or at '/tmp/Ultralytics/settings.json'
yolo_bdt | Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see https://docs.ultralytics.com/quickstart/#ultralytics-settings.
yolo_bdt | INFO: Started server process [1]
yolo_bdt | INFO: Waiting for application startup.
yolo_bdt | INFO: Application startup complete.
yolo_bdt | INFO: Uvicorn running on http://0.0.0.0:8080 (Press CTRL+C to quit)

```


Image 3.2 (Docker Desktop): Displays the container yolo_bdt in a **RUNNING** state. The logs confirm that the **Uvicorn** server is active on `http://0.0.0.0:8080`, mapped to the host port 8080.



4 Accuracy Discussion and Performance Reflection

The custom YOLO model demonstrated high precision on high-contrast objects, achieving over 93% confidence on both the 5 Taka coin and the 1000 Taka banknotes. While the system showed excellent multi-object detection capabilities in the 1000 Taka test case, performance fluctuated on smaller denominations with lower confidence scores (averaging 35%) due to grayscale noise and complex backgrounds. Notably, the 20 Taka test resulted in a false negative, suggesting that the model's current weights are sensitive to banknote orientation and specific reverse-side features. Overall, the system proved highly reliable for flat, clear denominations but would benefit from further data augmentation to improve its sensitivity to lower-value notes in varied lighting.

Module 5: Dockerization, Deployment, and API Interaction

5.0 Installing Docker

Before building your YOLO image, you must have the Docker Engine installed on your host machine.

For Windows and macOS

1. **Download:** Go to the official [Docker Desktop](#) website.
2. **Install:** Run the installer and follow the wizard.
3. **Windows Note:** Ensure **WSL 2 (Windows Subsystem for Linux)** is enabled when prompted for better performance.
4. **Start:** Launch Docker Desktop and wait for the "Engine Running" icon to turn green.

For Linux (Ubuntu Example)

Run the following commands in your terminal to install the Docker engine:

Bash

Update packages

```
sudo apt-get update
```

Install Docker

```
sudo apt-get install docker.io -y
```

Start and enable Docker service

```
sudo systemctl start docker
```

```
sudo systemctl enable docker
```

Allow your user to run docker without 'sudo' (Optional but recommended)

```
sudo usermod -aG docker $USER
```

5.1 Essential Docker Lifecycle Commands

Use these commands to manage your YOLO detection service.

A. Image Management

Bash

List all locally stored images

```
docker images
```

Remove an old image to save space

```
docker rmi yolo-currency-app
```

B. Container Operations

Bash

1. Build the image from your Dockerfile

```
docker build -t yolo-currency-app .
```

2. Run the container (Mapping Port 8080)

-d: Detached mode (background)

-p: Map host port 8080 to container port 8080

--name: Assign a readable name to the container

```
docker run -d -p 8080:8080 --name bdt-detector yolo-currency-app
```

3. Check if the container is running

```
docker ps
```

4. Stop/Restart/Remove the container

```
docker stop bdt-detector
```

```
docker start bdt-detector
```

```
docker rm bdt-detector
```

C. Debugging and Monitoring

Bash

View real-time logs from the Python server

```
docker logs -f bdt-detector
```

Access the container's internal terminal for advanced debugging

```
docker exec -it bdt-detector /bin/bash
```

5.2 Using the API Endpoint

The API accepts image files and returns structured JSON data containing denomination and confidence levels.

Endpoint Details

- **URL:** `http://127.0.0.1:8080/predict`
- **Method:** `POST`
- **Content-Type:** `multipart/form-data`

Interaction via Postman

1. Set the request type to **POST**.
2. Enter the URL: `http://127.0.0.1:8080/predict`.
3. In the **Body** tab, select **form-data**.
4. Add a key named **"file"**, change the type to **File**, and upload your currency image.
5. Click **Send**.

Interaction via cURL

Bash

```
curl -X POST http://127.0.0.1:8080/predict -F "file=@/path/to/your/image.jpg"
```

Expected JSON Response

JSON

```
{  
  "detected_classes": ["1000"],  
  "confidence_scores": [0.9321],  
  "coordinates": [[0.0, 320.03, 626.27, 638.40]]  
}
```

5.3 Final Integration Workflow

To deploy your project from scratch:

1. **Install Docker** (Step 5.0).
2. **Navigate** to your project directory (containing `Dockerfile`, `main.py`, and `custom_weights.pt`).
3. **Build:** `docker build -t yolo-currency-app .`
4. **Run:** `docker run -d -p 8080:8080 --name bdt-detector yolo-currency-app`
5. **Test:** Use Postman or cURL to hit the endpoint with a BDT note or coin image.